

codsoft-ml-2

August 5, 2024

TASK-2 CUSTOMER CHURN PREDICTION

```
[1]: import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, \
    accuracy_score
```

```
[2]: #Load data
df = pd.read_csv('/content/WA_Fn-UseC_-Telco-Customer-Churn.csv')
```

```
[3]: df.head()
```

```
[3]:   customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0  7590-VHVEG  Female                0      Yes           No         1           No
1  5575-GNVDE   Male                0      No            No        34           Yes
2  3668-QPYBK   Male                0      No            No         2           Yes
3  7795-CFOCW   Male                0      No            No        45           No
4  9237-HQITU  Female                0      No            No         2           Yes
```

```
      MultipleLines  InternetService  OnlineSecurity  ...  DeviceProtection  \
0  No phone service              DSL                No  ...                No
1                No              DSL                Yes  ...                Yes
2                No              DSL                Yes  ...                No
3  No phone service              DSL                Yes  ...                Yes
4                No      Fiber optic                No  ...                No
```

```
      TechSupport  StreamingTV  StreamingMovies  ...  Contract  PaperlessBilling  \
0                No           No              No  ...  Month-to-month           Yes
1                No           No              No  ...    One year             No
2                No           No              No  ...  Month-to-month           Yes
3                Yes           No              No  ...    One year             No
4                No           No              No  ...  Month-to-month           Yes
```

```
      PaymentMethod  MonthlyCharges  TotalCharges  Churn
0      Electronic check           29.85          29.85   No
```

1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

[5 rows x 21 columns]

```
[4]: print(df.columns)
```

```
Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
      'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
      'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')
```

```
[5]: df.shape
```

```
[5]: (7043, 21)
```

```
[6]: # Data preprocessing
      # Encode categorical variables with One-Hot Encoding
      df.replace(' ', pd.NA, inplace=True)
      df.dropna(inplace=True)

      cat_cols = ['gender', 'Partner', 'Dependents', 'PhoneService',
                  'MultipleLines', 'InternetService', 'OnlineSecurity',
                  'OnlineBackup', 'DeviceProtection', 'TechSupport',
                  'StreamingTV', 'StreamingMovies', 'Contract',
                  'PaperlessBilling', 'PaymentMethod']

      df_encoded = pd.get_dummies(df, columns=cat_cols, drop_first=True)
```

```
[7]: # Split data into features (X) and target (y)
      X = df_encoded.drop(['customerID', 'Churn'], axis=1) # Assuming 'Churn' is the
      ↪ target variable
      y = df_encoded['Churn']
```

```
[8]: # Split data into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=28)
```

```
[9]: # Standardize features
      scaler = StandardScaler()
      X_train = scaler.fit_transform(X_train)
      X_test = scaler.transform(X_test)
```

```
[10]: # Model 1: Logistic Regression
log_reg = LogisticRegression(max_iter=1000, random_state=28)
log_reg.fit(X_train, y_train)

# Predictions and Evaluation
y_pred_log_reg = log_reg.predict(X_test)
print("Logistic Regression:")
print(classification_report(y_test, y_pred_log_reg))
print(confusion_matrix(y_test, y_pred_log_reg))
print(accuracy_score(y_test, y_pred_log_reg))
```

Logistic Regression:

	precision	recall	f1-score	support
No	0.86	0.90	0.88	1020
Yes	0.70	0.60	0.65	387
accuracy			0.82	1407
macro avg	0.78	0.75	0.76	1407
weighted avg	0.81	0.82	0.82	1407

```
[[920 100]
 [154 233]]
0.8194740582800284
```

```
[11]: # Model 2: Random Forest Classifier
rf = RandomForestClassifier(n_estimators=50, random_state=28)
rf.fit(X_train, y_train)
```

```
[11]: RandomForestClassifier(n_estimators=50, random_state=28)
```

```
[12]: # Predictions and Evaluation
y_pred_rf = rf.predict(X_test)
print("\nRandom Forest Classifier:")
print(confusion_matrix(y_test, y_pred_rf))
print(classification_report(y_test, y_pred_rf))
print("Accuracy:", accuracy_score(y_test, y_pred_rf))
```

Random Forest Classifier:

```
[[930 90]
 [200 187]]
```

	precision	recall	f1-score	support
No	0.82	0.91	0.87	1020
Yes	0.68	0.48	0.56	387
accuracy			0.79	1407

macro avg	0.75	0.70	0.71	1407
weighted avg	0.78	0.79	0.78	1407

Accuracy: 0.7938877043354655